

AW ROSTAMANI GROUP IT

Vehicle Live Tracking

Architecture presentation · Product Engineering Manager assignment

Hitesh Pachpor · 12 September 2026

The scenario

AW Rostamani (AWR) works with third-party logistics (3PL) vendors who pick up and drop off customers' vehicles, for example between a showroom and a customer's home.


Once a trip has started, the AWR operations team needs to see where the vehicle is on a map in a web dashboard, and the map should update without refreshing the page.

Each trip involves these roles:

- **AWR Operations** creates trips and tracks them.
- **Vendor Controller** works for the 3PL vendor and assigns each trip to one of the vendor's drivers.
- **Vendor Driver** starts the trip, shares their location while driving, and ends the trip.

The demo data is based in the UAE and uses AWR's automotive brands (Nissan, INFiniti, Renault, Chery and Zeekr).

The scenario



What the demo version includes

Requirement from the brief	How it is covered in the demo version
Trip management API	Trips can be created, listed, viewed, assigned to a driver, started, completed and cancelled.
Location ingestion endpoint	<code>POST /api/trips/:id/location</code> accepts latitude, longitude, timestamp, and an optional speed and event ID.
Live updates without polling	Server-Sent Events (SSE). Updates missed while disconnected are sent again on reconnect.
Trip simulator	Runs on the server, follows the driving route from Mapbox, and sends pings through the same location API as a real driver.
Data layer	PostgreSQL 17.
Active trips list	Trip list with status, search and pagination. It updates live when trips change.
Live map view	Mapbox map with a vehicle marker that moves as each ping arrives.
Trip timeline	A log of location pings showing device time, received time, speed and source.
Vendor simulator panel	The driver can start and stop a simulation and set the update interval and distance per ping.
Responsive layout	Works on desktop, tablet and mobile.

What the demo version includes

The demo version also includes the following, which the brief did not ask for:

- A separate workspace for each role.
- Rules that stop a driver or a vehicle from being on two live trips at the same time.
- Real GPS from the driver's browser, with an offline queue that holds pings until the network is back.
- Pasting a Google Maps link to fill in the pickup and drop-off locations when creating a trip.
- Docker Compose to run the whole stack with one command, and UAE-based seed data.

Login in the demo version is a simple role switcher, and the APIs are not protected. The brief allows authentication to be covered in the presentation only, and the proposed design is in section 6.

1.1 Vendor driver flow



1.1 Vendor driver flow

Live trip

The driver allows location access in the browser and starts the trip. The browser reads the location every 5 seconds and sends it to the server.


If the network drops, pings are kept in a queue on the device (up to 720 pings) and sent once the connection is back. Each ping keeps its original time and a unique event ID, so a ping that gets sent twice is only saved once.

Simulated trip

The driver chooses how often a ping is sent (1 to 60 seconds) and how far the vehicle moves between pings (0.1 to 20 km). The server starts the trip and moves the vehicle along the driving route shown on the map. Each ping goes through the same location API that a real driver uses.

After the last ping at the drop-off point, the trip is completed automatically. If the simulation is stopped early, the trip stays in transit.

1.1 Vendor driver flow



1.1 Vendor driver flow


Trip rules

- A driver must be assigned before the trip can start.
- Only the assigned driver can start or complete a trip. Operations can cancel it.
- A driver can have only one trip in transit at a time.
- A vehicle can be on only one scheduled or in-transit trip at a time.
- A driver's scheduled trips must be at least three hours apart.
- If two people update the same trip at the same time, the second update is rejected so it does not overwrite the first.

Timestamps on each ping

Each ping stores the time the device recorded it and the time the server received it. For example, if a driver loses signal in a basement car park, operations can see that the location was recorded at 14:01 and arrived at 14:04.

1.2 AWR operations team flow



1.2 AWR operations team flow

Monitor the dashboard

- See all trips, search them, and filter by Scheduled, In transit, Completed or Cancelled.
- The list updates by itself when a trip is created, assigned, started, completed or cancelled.
- Create a trip by choosing a customer, vehicle, vendor, pickup location and drop-off location. A Google Maps link can be pasted in place of an address. The driver is chosen later by the vendor.

View trip detail

- Watch the vehicle move on the map, with the last 100 location pings listed beside it.
- See whether the live connection is active.
- Cancel a trip that is scheduled or in transit.
- On a completed trip, the vehicle's last position is labelled "Actual drop-off" when it is 50 metres or more from the planned drop-off point.

Starting, simulating and completing a trip are left to the driver, so operations does not see those buttons.

1.2 AWR operations team flow

Alerts (proposed, not in the demo version)

Alert	When it is raised	Who is notified
Location has gone quiet	No ping for a few minutes while the trip is in transit	Operations, so they can call the vendor
Trip not started	The scheduled time has passed and the trip has not started	Operations and the vendor controller
Off route	The vehicle is too far from the planned route	Operations
Arrived at or left a zone	The vehicle enters or leaves the pickup point, drop-off point or a yard	Operations
Drop-off mismatch	The trip ended far from the planned drop-off point	Operations
Trip started or completed	The trip status changes	The customer (by SMS) and the CRM

The first alert to build would be "location has gone quiet". Without it, operations only finds out about a problem when someone notices the marker has stopped moving.


1.3 Admin flow

The demo version has no separate admin role. In a 3PL setup, AWR creates the trip and the vendor assigns the driver, so the admin tasks from the brief are split between the existing roles.

Admin task	Who does it in the demo version	Proposed future version
Create trips	AWR Operations	Stays with AWR Operations
Assign trips to drivers	Vendor Controller	Add rules for reassigning trips and meeting SLA time windows
View history	Trip list filters and the location ping log	Longer retention, export and trip replay
Manage customers, vehicles, vendors and drivers	Loaded from seed data and read-only	Screens to add and edit these records
Manage user access	Demo role switcher	Role management, vendor onboarding and an audit log

Later on, an admin role would also manage vendor API keys, SLA and geofence settings, and the audit log of who viewed a live trip.

2.1 High-level architecture diagram



2.1 High-level architecture diagram

Technology stack

Next.js 16 (App Router), Zod for request validation, Drizzle ORM, PostgreSQL 17, Mapbox for maps and driving routes, Docker Compose, and Vitest with Testcontainers for tests.

Long-running Node process

The live connections to the dashboard and the simulator's timer both need a process that keeps running. Serverless functions stop between requests, which would close those connections and stop running simulations.

Choice of data store

The brief suggested in-memory storage, Redis or a lightweight database.

- **In-memory storage** is quick to set up, but trip history is lost when the process restarts, and missed updates cannot be sent again after a reconnect.
- **Redis** handles messaging well, but it is less suited to storing trip records and a long history of location pings.
- **PostgreSQL** was chosen. It stores trips and the full ping history, and its unique indexes enforce the trip rules. Its built-in notification feature (`LISTEN` and `NOTIFY`) is used to tell the app when a new ping is saved, so one database covers both needs while the app runs as a single instance.

2.2 Justification for real-time mechanism chosen

	Polling	WebSockets	Server-Sent Events (SSE)
Direction	The dashboard repeatedly asks the server for updates	Both ways	Server to dashboard only
Fit for this dashboard	Adds load, and updates arrive late	Adds a two-way channel the dashboard does not use	Suits a dashboard that only receives updates
Reconnecting	Has to be built	Has to be built	Built into the browser
Support in Next.js	Simple	Needs an extra library and server setup	Works with a standard streaming response


SSE was chosen because the dashboard only receives location updates and never sends data back over that connection.

If the connection drops, the browser reconnects by itself and sends the ID of the last update it received. The server then sends every ping saved after that ID, so the trail on the map has no gaps.

The server also sends a small heartbeat every 15 seconds, so proxies and load balancers do not close the connection as idle.

WebSockets would be worth adding if the system later needs to send messages to drivers during a trip.

2.3 How location data flows from the vendor interface to the AWR dashboard



2.3 How location data flows from the vendor interface to the AWR dashboard

Saving a ping

1. The driver's browser or the simulator sends the location to `POST /api/trips/:id/location`.
2. The server checks the latitude, longitude, timestamp and optional speed.
3. The ping is rejected if the trip is not in transit.
4. The ping is saved. If a ping with the same event ID already exists for this trip, it is not saved again and no notification is sent.

Showing it on the dashboard

1. When a trip is opened, the dashboard connects to `GET /api/trips/:id/events`.
2. If the dashboard is reconnecting, the pings it missed are sent first (up to 1,000).
3. After that, each new ping is pushed as soon as it is saved.
4. The map marker moves to the new position, and the ping log shows the latest 100 entries.

Pings from real drivers and from the simulator follow the same path. They are told apart by a `source` field set to `vendor` or `simulator`.

2.4 How you would handle scale (100+ concurrent trips)


Current capacity of the demo version

- 100 trips sending a ping every 5 seconds adds up to about 20 saves per second, or about 1.7 million rows per day. PostgreSQL handles this load comfortably with an index on trip and time.
- 20 operations users each watching one trip means 20 open connections, which is a small load for one Node process.
- The dashboard opens a live connection only for the trip being viewed, so a list of 100 trips does not open 100 connections.

Limits with a second app instance

- The database listener runs inside one process. Users connected to another instance would not get live updates until they reconnect.
- Running simulations are held in the memory of the process that started them.

2.4 How you would handle scale (100+ concurrent trips)



2.4 How you would handle scale (100+ concurrent trips)

Proposed changes for more than one instance

- PostgreSQL continues to store every ping.
- A message bus such as Google Pub/Sub or Redis takes over from the database notification. Every instance receives every new ping and pushes it to its own connected users.
- The simulator moves to a separate worker, so each simulation runs in one place.
- Old location pings are archived once the history covers many months.

Kafka would be more than this volume needs. Until a second instance is required, a single instance with a health check is enough.

3.1 API contract design

Method and path	What it does
POST /api/trips	Create a trip
GET /api/trips	List trips, filtered by status, vendor or driver
GET /api/trips/:id	Get a single trip and its current status
PATCH /api/trips/:id	Assign a driver, or change the status (start, complete, cancel)
POST /api/trips/:id/location	Send a location ping
GET /api/trips/:id/positions	Get the recent location pings for a trip
GET /api/trips/:id/events	Live stream of new pings for one trip
GET /api/trips/events	Live stream of changes to the trip list
GET, POST, DELETE /api/trips/:id/simulation	Check, start or stop a simulation

There are also a few supporting endpoints. They return the customers, vehicles, vendors and drivers used in the forms, convert a Google Maps link into coordinates, and report whether the app can reach the database (health check).

3.1 API contract design

Request example: send a location ping

```
{
  "lat": 25.2048,
  "lng": 55.2708,
  "timestamp": "2026-09-10T18:00:00Z",
  "speed": 12.5,
  "eventId": "vendor-message-123"
}
```

Error response format


All errors use the same format. Error codes stay the same between releases, so a vendor's app can rely on them.

```
{ "error": { "code": "DRIVER_NOT_ASSIGNED", "message": "A driver must be assigned before the trip can start.", "details": {} } }
```

Versioning strategy (proposed)

- Vendor-facing APIs are published through Apigee under `/api/v1`.
- Within a version, fields can be added but are never removed or renamed.
- A breaking change is released as `/api/v2`, and both versions run side by side while vendors move over.
- The OpenAPI documentation is generated from the existing Zod schemas, so it stays in line with the code.

3.2 Webhook or event-driven approach for notifying downstream systems



3.2 Webhook or event-driven approach for notifying downstream systems

The API publishes an event whenever something happens to a trip, and each downstream system subscribes to the events it needs. If the API called the SMS provider directly, a slow or failed SMS request could delay or break the trip start.

Event	When it is published	Who would use it
trip.created	Operations creates a trip	Vendor systems
trip.driver_assigned	The controller assigns a driver	Driver app notification
trip.started	The trip goes in transit	Customer SMS, CRM
trip.position	A sample of pings, not every ping	Analytics
trip.stale , trip.deviated	The alert service detects a problem	Operations, vendor SLA reports
trip.completed , trip.cancelled	The trip ends	Customer SMS, ERP job closure, billing

External systems that cannot subscribe to the bus receive the same events as webhooks. Each webhook is signed so the receiver can check that it came from AWR, and failed deliveries are retried.

3.3 Middleware layer considerations

Already in the demo version

- Requests that change data are validated with Zod.
- All errors use one format.
- When a request breaks a database rule, the API returns a specific error (409 Conflict or 422 Unprocessable) with a clear code, in place of a generic server error.

Proposed in the app

Area	Approach
Authentication	Check the user's token or the device's API key on every request
Request validation	Keep Zod validation on every request
Duplicate requests	Location pings already use an event ID. Trip creation would accept an Idempotency-Key header
Logging	JSON logs with a request ID, trip ID and vendor ID. Exact coordinates are left out of normal logs
Rate limiting	Limit each vendor key on the location API, for example to one ping per second with a small burst allowance

3.3 Middleware layer considerations

Proposed at the API gateway (Apigee)

HTTPS, API key checks, usage quotas, separate API products for vendors and AWR staff, and reports on failed requests.

A rate limit inside the app works only while there is one instance. Once there are more, Apigee enforces the limit and the app keeps a basic check as a fallback.


3.4 How the system could integrate with AWR's existing operational platforms

AWR's internal system names were not available, so these assumptions were made from public information:

- Customer journeys are managed in a CRM and a dealer management system.
- Vendor jobs are closed and costed in an ERP.
- APIs are published through Apigee.
- Group data is stored in BigQuery.

Platform	How the tracking system would connect to it
Apigee	Publishes the location API to vendors, with keys, quotas and HTTPS
CRM or dealer management system	Creates trips automatically, for example for a service pickup or a new car delivery
ERP	Receives <code>trip.completed</code> to close and cost the vendor job
Customer notifications service	Receives <code>trip.started</code> and <code>trip.completed</code> to send SMS or WhatsApp messages
BigQuery	Receives trip and location data for SLA reporting
Vendor systems	Send GPS pings to the location API

4.1 Moving to an event-driven architecture



4.1 Moving to an event-driven architecture

After a ping is saved, the location API publishes a `location.recorded` event to a message queue. The dashboard, the alert service, BigQuery and the webhook sender each read from the queue separately.

With this in place, a slow or failing consumer has no effect on the location API, and new consumers can be added without changing it. Every app instance also receives every update, which is needed once there is more than one instance.

This change is worth making when a second app instance or a second consumer of location data is added.

4.2 Geofencing

Geofencing raises an alert when a vehicle enters or leaves a defined area.

- Each trip gets a circular zone around the pickup and drop-off points, for example with a 150-metre radius. Yards and service centres can have custom shapes.
- A separate service reads new pings from the queue and checks them against the trip's zones, so the location API does not slow down.
- The service raises events such as `trip.arrived_pickup`, `trip.left_pickup` and `trip.arrived_dropoff`.
- Circular zones only need a distance calculation. PostGIS can be added when custom shapes are introduced.

The same zones could be used to complete a trip automatically at the drop-off point and to send the customer an arrival message.

4.3 Historical trip replay and route deviation detection

Trip replay

Every ping is already stored in order, together with the planned route. Replay would add a timeline slider to a completed trip that moves the marker through the stored pings. Operations could use it when looking into a customer complaint or a dispute with a vendor.

Route deviation detection

The system measures the distance between each ping and the planned route. If the vehicle stays further away than a set limit (for example 500 metres) for longer than a set time, a `trip.deviated` alert is raised.

After the trip, the same measurement can show how closely each vendor's drivers follow planned routes. The demo version already does a basic form of this check when it labels the actual drop-off point.

4.4 Multi-region or edge deployment for lower latency

The users are in the UAE, so the biggest improvement comes from running the app and database in a Google Cloud region close to the UAE.

- Static files such as JavaScript, CSS and images are served through Cloud CDN. Mapbox serves map tiles from its own global network.
- A single primary database with a read replica is enough. Users are in one country, so writing to databases in several regions is not needed.
- If AWR operates this system in other countries later, each country can have its own deployment, which also keeps location data inside that country.

5.1 Map rendering performance with frequent marker updates

In the demo version

- The map is created once and reused.
- The driving route is fetched once per trip, and the map zooms to fit the route once, not on every ping.
- The ping log shows at most the latest 100 entries.
- Map animations are turned off for users who have reduced motion enabled.

Proposed improvement

The map markers are currently removed and drawn again each time the trip data changes. At one ping every 5 seconds this is not noticeable. For faster updates, the vehicle marker would be moved to its new position with a short animation, so it glides between points.

5.2 Batching vs streaming location payloads

Drivers send pings one at a time as they are recorded, so operations sees the latest position straight away. Pings are sent in a batch only after the device has been offline, when the queued pings are sent together.

The server streams each saved ping to the dashboard as soon as it is saved. If a vendor sent a ping every second, the dashboard could receive a reduced stream, for example one update every 2 to 3 seconds, while every ping is still stored in the database.

SMS, CRM and analytics systems receive trip events and a sample of pings, not every ping.

5.3 CDN and asset caching strategy for the web dashboard

Content	Caching approach
JavaScript, CSS and fonts	Served through Cloud CDN and cached for a long time. File names change with each release, so users get new files after a deployment
Map tiles	Served and cached by Mapbox's own network
Pages and API responses	Not cached, because they contain live and user-specific data
Live updates (SSE)	Never cached. The <code>X-Accel-Buffering: no</code> header stops proxies from holding updates back

The Mapbox access token can be seen in the browser, so it is restricted to AWR's domains.

5.4 Database indexing strategy for location records

Indexes in the demo version

- Location pings are indexed by trip and time, so a trip's timeline and latest position load quickly.
- Trips are indexed by status and last update time, so the trip list loads quickly.
- Partial unique indexes enforce one live trip per driver and one active trip per vehicle.
- Event IDs are unique per trip, so the database rejects duplicate pings.

Proposed as data grows

- Split the location table by month, so older months can be archived or deleted easily.
- Add PostGIS spatial indexes when geofencing with custom shapes is introduced.

5.5 Connection management for WebSocket at scale

The demo version uses SSE, and the same concerns apply.

- Each dashboard opens one live connection for the trip being viewed and one for the trip list.
- A heartbeat every 15 seconds stops proxies from closing connections that look idle.
- The browser reconnects by itself after a drop, and missed pings are loaded from the database.
- When a user closes the page, the server stops sending to that connection and removes it.
- At larger scale, HTTP/2 lets many streams share one connection, and a message bus lets any app instance serve any user (see 2.4).

5.6 Server-side rendering vs client-side rendering trade-offs in Next.js

In the demo version

The pages are light server components, and data is loaded in the browser. The demo login is stored in the browser, so the server cannot tell who the user is when it renders a page.

Proposed once real login is in place

Part of the page	Where it is rendered	Reason
Trip header (vehicle, customer, status, driver)	Server	Shows useful content as soon as the page loads
First page of the trip list	Server	Shows useful content as soon as the page loads
Live map, ping log and live connection	Browser	Changes every few seconds and relies on browser features

The trip details appear first, and the location history loads and streams in afterwards.

6.1 Authentication (who are you?)

Each type of user signs in differently.

User	Who they are	Proposed sign-in method
AWR Operations and Admin	AWR employees	Sign in through AWR's company identity provider using OAuth2 (Authorization Code with PKCE). The session is kept in a secure, HTTP-only cookie
Vendor Controller	Vendor staff	Same sign-in method, with each vendor set up separately. The vendor ID is included in the user's token
Driver's device	A phone or vendor system sending GPS	A short-lived device token or vendor API key, limited to that driver and their trips. Tokens can be rotated

Device keys

A device only needs permission to send pings for its own trips, so it gets its own key and does not use the dashboard login. If a phone is lost, its key can be revoked without affecting anyone's dashboard access.

Live connections

The browser's built-in SSE client cannot add custom headers, so the live connection uses the same secure cookie as the rest of the dashboard. Tokens are kept out of URLs, where they would end up in logs.

Simulator

The simulator is turned off in production, or limited to admins with each use recorded.

6.2 Authorization (what can you do?)

The server checks every request against the user's role and the vendor or driver named in their token.

Action	Vendor Driver	Vendor Controller	AWR Operations	Admin
Create a trip	No	No	Yes	Yes
Assign a driver	No	Own vendor's trips	View only	Yes, recorded in the audit log
Start, complete or simulate	Own assigned trips	No	No	Emergency use only
Send location	Own in-transit trip	No	No	No
Cancel a trip	No	No	Yes	Yes
View trips and live map	Own trips	Own vendor's trips	All	All
Manage master data, API keys and SLAs	No	View own drivers	View	Yes
Replay and export history	No	Own vendor's trips	Yes	Yes

The demo version applies these rules in the browser only. In production the server runs the same checks. A request without a valid login gets 401 Unauthorized, and a request for another vendor's or driver's trip gets 403 Forbidden.

A vehicle's live location is sensitive, so each time someone views it, the view is recorded in an audit log.

6.3 Data in transit

- All traffic to the dashboard and APIs uses HTTPS (TLS 1.2 or higher), and HTTP requests are redirected to HTTPS.
- The SSE stream runs over the same HTTPS connection. If WebSockets are added for drivers later, they use secure WebSockets (WSS).
- The app connects to the database over an encrypted connection inside a private network.
- Tokens are never put in URLs, so they do not show up in server or proxy logs.

6.4 Data at rest

- Cloud SQL encrypts the database, its backups and its replicas. AWR can manage the encryption keys if required.
- Coordinates are not encrypted field by field, because the system needs to query them to draw maps and calculate distances. Access to the database is restricted instead.
- Only the app's service account can read location data. Staff access to the database is limited and logged.
- Detailed location history is kept for a fixed period, for example 90 days, and then deleted or kept only as summaries.
- Customer contact details have the same access controls, and exact coordinates are left out of normal application logs.

6.5 API abuse prevention

The location API receives traffic from outside AWR, which makes it the most exposed part of the system.

Control	What it protects against
Input validation	Invalid coordinates, timestamps and speeds are rejected before they reach the database (in the demo version)
Unique event IDs	A retried ping is not saved twice (in the demo version)
Trip must be in transit	Pings cannot be sent to a trip that has not started or has already ended (in the demo version)
One live trip per driver	The same driver cannot be active on two trips (in the demo version)
Rate limiting per device	A device or key cannot flood the API. About one ping per second is plenty for a vehicle
Anomaly detection	Flags impossible speeds, sudden jumps in location, and repeated failed logins
Device tokens and gateway quotas	Only registered devices can send pings

7.1 Operational metrics

Metric	How it is calculated
Trip duration	Time from trip start to trip completion, compared with the planned time
Average speed	Average and 95th percentile speed from the location pings
Delay frequency	Share of trips that started late or arrived later than the route estimate
Route efficiency	Distance actually driven compared with the planned route distance
Tracking quality	Number of trips where the location went quiet or the drop-off point did not match

7.2 System health

Metric	What it shows
API response time (95th percentile)	Whether the location API is slowing down
Error rate by error code	Whether vendors are sending bad data or the system is failing
Ping delay (time received minus time recorded)	Whether drivers are losing signal or pings are arriving late
Live connection stability	Number of open connections, reconnects and dropped connections
Uptime	Results of the health check endpoint

7.3 Business metrics

Metric	What it shows
SLA compliance per vendor	Share of trips that met the agreed pickup and drop-off times
On-time pickup and drop-off rate	How reliable each vendor is
Route deviation rate per vendor	How often drivers leave the planned route
Trip volume by vendor, brand and emirate	Where more vendor capacity is needed
Cost per completed trip	Vendor cost, once the ERP is connected

7.4 Tools

Purpose	Tool	Reason
Data warehouse	BigQuery	AWR's group data is already stored there
Loading data into the warehouse	Pub/Sub and Dataflow	Streams trip events into BigQuery without adding load to the app
Operational and business dashboards	Looker or Looker Studio	Vendor SLA scorecards and operations reports
System health	Google Cloud Monitoring and Logging	Response times, errors and dropped connections, with alerts
Product usage	Google Analytics 4 (optional)	How the dashboard is used, for example how long it takes to create and assign a trip

Mixpanel or a similar product analytics tool is not needed at the start. It can be added later if AWR already uses it elsewhere.

Approach by layer

Layer	Approach	Covered in the demo version	Proposed next
Unit tests	Core business logic	Trip status changes, driver availability rules, request validation, route and distance calculations, the offline ping queue, and role access	Add tests with each new feature
Integration tests	API routes and the live connection lifecycle	API handler tests, and tests against a real PostgreSQL database (using Testcontainers) for database rules, duplicate pings and database notifications	HTTP tests against a running server
Component tests	UI components with React Testing Library	Not covered yet	Trip list, driver assignment and live connection status
E2E tests	Playwright: trip simulation and live map update	Not covered yet	Sign in, assign a driver, run a simulation, and check that a new ping moves the map marker
Performance tests	Load test of the location API (k6 or Artillery)	Not covered yet	Around 20 pings per second with 50 open live connections, run nightly

Approach by layer

Focus of the current tests

Most of the current tests cover issues that are hard to spot by using the app, such as a driver being put on two trips at once or a ping being saved twice after a retry.

When each type of test runs

- Every pull request runs lint, type checks, unit tests and a production build. These commands already exist in the project.
- Integration tests need Docker, so the CI pipeline runs them in an environment where Docker is available.
- End-to-end tests run against a preview environment before changes reach staging.
- Performance tests run nightly against staging, so they do not slow down pull requests.


9.1 Hosting platform for Next.js

AWR already runs digital products on Google Cloud, so the proposal is to host this system there too.


Option	Assessment
Vercel	Very good Next.js support. Its functions are short-lived, though, so long-running live connections, the database listener and the simulator would need to run somewhere else. It would also be one more platform to manage
AWS	Would work technically, but adds cost and effort if AWR's other platforms are on Google Cloud
Google Cloud (GKE and Cloud SQL)	Supports long-running processes and uses platforms AWR already has

Cloud Run could also work, as long as at least one instance is always running and the simulator runs in a separate worker. Without those two changes, idle instances would shut down and drop live connections.

9.1 Hosting platform for Next.js



9.2 CI/CD pipeline design



9.2 CI/CD pipeline design

Branch strategy

Feature branches are kept short and merged into `main` through pull requests. `main` is kept ready to deploy at all times, and production releases are tagged.

Automated tests

Every pull request runs lint, type checks, unit and integration tests, and a production build. A Playwright test on a preview environment and a Terraform plan can be added later.

Staged rollout

Each release goes to dev first, then staging, then production. Production deployments need approval and are rolled out gradually. New features, such as extra location fields, are released behind feature flags so they can be turned off without a new deployment.

With a single instance, a deployment briefly drops live connections. Browsers reconnect by themselves and load the pings they missed, so no data is lost. With two or more instances, a rolling update avoids the drop.

9.3 Environment management

Environment	Data	Used by
Local	Docker Compose with seed data	Engineers
Dev	Shared Google Cloud project with test data	Engineers, for integration work
Staging	Close to production, with anonymised data and no real customer locations	AWR IT and operations, for user acceptance testing
Production	Real vendors and trips, stricter access, simulator turned off	Restricted access

- Each environment has its own Google Cloud project, database and API keys.
- Secrets are stored in Secret Manager, not in code or Docker files.
- Seed data and demo passwords are used only locally and in dev.

9.4 Monitoring & alerting

Signal	Starting threshold	Action when exceeded
Uptime (health check)	99.9%	Alert the on-call engineer
Location API response time (95th percentile)	Under 200 ms	Alert if it stays above the threshold for 5 minutes
Location API error rate	Under 0.1%	Alert the on-call engineer
Live connection drop rate	Alert if reconnects rise sharply	Check proxies and load balancers for closed or buffered connections
Trips with no recent location	Tracked as a product metric	Alert operations and check the vendor and the location API
Conflicting trip updates	Expected to be rare	Investigate a sudden increase, as it usually points to a bug

Google Cloud Monitoring and Logging cover the platform, and an error tracking tool covers the Next.js app. The `/api/health` endpoint is used for both the uptime check and the Kubernetes health check.

9.5 Infrastructure as Code considerations

All infrastructure is defined in Terraform and reviewed through pull requests in the same way as application code. This includes:

- GKE cluster and app deployment
- Cloud SQL database and backups
- Access roles (IAM) and service accounts
- Secret Manager
- Apigee API products
- Cloud CDN and load balancer
- Pub/Sub topics
- Monitoring alerts and dashboards

Dev, staging and production are then created from the same definitions, with no manual setup in the console.

Terraform is the default choice because it is widely used and well supported on Google Cloud. Pulumi would also be fine if the team already uses it.

What would be done differently with more time

In the demo version	With more time	Reason it was deferred
Demo role switcher and open APIs	Real sign-in, device tokens and server-side access checks	The brief allows authentication to be covered in the presentation only
Database notifications within one process	Message bus (Pub/Sub or Redis)	One instance was enough to run the full flow end to end
Simulations held in the app's memory	A separate simulator worker	The simulator is a tool for demos and testing
Unit and integration tests	Component, Playwright and k6 tests	The data rules and live update logic were tested first
No downstream notifications	<code>trip.started</code> and <code>trip.completed</code> events and webhooks	There were no real SMS or CRM systems to connect to

What would be done differently with more time

In the demo version	With more time	Reason it was deferred
Data loaded in the browser	Trip details rendered on the server	Needs real sign-in first
API without a version number	<code>/api/v1</code> behind Apigee, with OpenAPI documentation	No external vendors were using the API yet
Map markers redrawn on each update	Vehicle marker moved with a short animation	Not noticeable at one ping every 5 seconds
No alerts	"Location has gone quiet" and "trip not started" alerts	The live tracking flow was built first
No CI pipeline configured	CI pipeline running the existing lint, type check, test and build commands	Checks were run locally during development